

4.0

Last week we've covered: Quicksort alg. that has a not-so-good worst ~~quicksort~~ case behavior, but a good average case behavior that works well in practice.

Then, we've seen several divide and conquer algs:

Binary search, Karatsuba's integer multiplication alg. and Strassen's matrix multiplication alg.

Today we are going to see yet another sorting alg. heapsort that is based on a neat data structure called heap. We will introduce heaps and discuss some of their properties.

Then we will prove a lower bound on comparison-based sorting algorithms like the one we've seen so far.


After that we will see a different sorting alg. that beats this lower bound. (because it is not based on comparisons!)

Heaps

A heap is a data structure - a way to organize your data - that has the form of a nearly complete binary tree.

That is, given an array A , each node of the tree corresponds to an element of the array.

For example, the array $A = [16, 4, 10, 14, 7, 9, 3, 2, 8, 1]$ is represented by the tree



Note that the tree is completely filled in each level except, possibly, the last one, which ~~is~~ is filled from the left up to a point.

Sometimes we will have a heap that represents only part of the array so we will have two parameters: $A.length$ and $A.heapsize$. (it we wish to consider only the first $A.heapsize$ elements of A)

The root of the tree is always $A[1]$, and given an index i , its left child is $2i$ and its right child is $2i+1$. Its parent is $\lfloor i/2 \rfloor$. Thus:

Parent(i)

Left(i)

Right(i)

1. Return $\lfloor i/2 \rfloor$

1. Return $2i$

1. Return $2i+1$

A maxheap (minheap) is a heap in which the values of the nodes satisfy a heap property. In maxheap the property is that $A[i] \geq A[\text{Parent}(i)]$ and in minheap $A[i] \leq A[\text{Parent}(i)]$.

Note that in maxheap the largest element is always at the root. Maxheaps are useful for sorting algorithms and for various other applications.

Another important parameter of a heap/tree is the height of a node. This is the number of edges on the longest downwards path from the node to a leaf. In the example the height of 1 is 3, the height of 3 is 1 and the height of a leaf is 0.

The height of the root in an n -element array is $\Theta(\lg n)$.

We will see several procedures for obtaining and maintaining a maxheap.

Max heapify : a procedure running in $\Theta(\lg n)$ time that is important for creating and maintaining a max heap.

Build Maxheap: a linear time procedure for rearranging A to be a maxheap.

Heapsort : An $\Theta(n \lg n)$ time in-place sorting alg.

We will also see several procedures that allow us to insert and remove elements from a heap which are useful when dealing with priority queues.

4.3


Let us start with Max heapify. This procedure gets an array A and an index i with the assumption that the trees rooted in i 's children are maxheaps. It then let the value of $A[i]$ "float down" to the right location to make the tree rooted at i also a maxheap.

Max-heapify (A, i)


1. $l = \text{Left}(i)$
2. $r = \text{Right}(i)$
3. If $l \leq A.\text{heapsize}$ and $A[l] > A[i]$
4. $\text{largest} = l$
5. Else $\text{largest} = i$
6. If $r \leq A.\text{heapsize}$ and $A[r] > A[\text{largest}]$
7. $\text{largest} = r$
8. If $\text{largest} \neq i$
9. exchange $A[i]$ with $A[\text{largest}]$.
10. Max-heapify ($A, \text{largest}$)

The alg first find the child with the largest value (if it is not i), then replaces $A[i]$ with $A[\text{largest}]$ and continues to "fix" the relevant subtree.

In our example if $i=2$ then $\text{largest}=4$ and we get



then $\text{largest}=9$ and we get



h.4

Runningtime: The first 9 lines take $\Theta(1)$ time and then we run the alg. again on a subtree. Thus if we measure the complexity in terms of the height of i we will get

$$T(h) \leq T(h-1) + \Theta(1)$$

$$\text{Hence } T(h) = \Theta(h) = \cancel{\Theta(\log n)} \Theta(\log n)$$

Correctness: It is easy to prove (say by induction on height) that Max-heapify is correct.


Given Max-heapify we can now give an alg. for building a maxheap from any array.

Notice that in a tree like ours (nearly complete) the elements in locations $A[\lfloor \frac{n}{2} \rfloor + 1, \dots, n]$ are always leaves of the tree. So ^{each} they already form a maxheap.

Build-Maxheap(A)

1. $A.\text{heapsize} = A.\text{length}$
2. For $i = \lfloor A.\text{length}/2 \rfloor$ to 1
3. Max-heapify(A, i)

In our example we will start from 5, Max-heapify will do nothing, then same with 4, 3. When we get to 2 we will get (as we saw) then we go to 1 and stop.



4.5

To show why it works we use the following invariant:

At the start of each For iteration, each node $i+1, i+2, \dots, n$ is a root of a max heap.

Initialization: $\frac{n}{2}+1, \dots, n$ are all leaves so invariant holds.

Maintenance: Follows because Max-heapify works (and by invariant, Left(i), Right(i) are roots of max heaps).

Termination: We stop when i=0 at this point node i=1 is a root of a max heap.

Run time: trivial bound is $O(n \log n)$ as we have $\frac{n}{2}$ iterations each taking $O(i) = O(\log n)$.

Better analysis: when a node at level h is called, Max-heapify runs in time $O(h)$.

OK, so how many nodes are at height h ?

Note that each such node has at least $2^{h-1}-1$ nodes below it. (a full tree at height $h-1$ has $2^{h-1}-1$ nodes below the root).

Thus, at most $\frac{n}{2^{h-1}}$ such nodes.

$$\begin{aligned} \text{Therefore runtime} &\leq \sum_{h=1}^{\log n} O(h) \cdot \frac{n}{2^{h-1}} = n \cdot O\left(\sum_{h=1}^{\log n} \frac{h}{2^{h-1}}\right) \\ &= O(n) \end{aligned}$$

4.6

Now that we have an efficient way to construct maxheap we can give an efficient sorting alg.

The idea is that in a maxheap the largest element is in the root. Thus we can exchange it with $A[1]$, reduce the heapsize by 1, fix the maxheap (Max-heapify) and repeat.

Heapsort(A)

1. Build-Max-heap(A)
2. For $i = A.length$ to 2
3. exchange $A[i]$ with $A[1]$
4. $A.heapsize = A.heapsize - 1$
5. Max-heapify(A, 1)

Runtime: $n-1$ iterations each taking $O(\lg n)$. overall $O(n \lg n)$

Correctness: Can prove that before for loop, elements in $i \dots n$ are the $n-i$ largest elements, sorted and that

~~Right(A, Left(i)) are maxheaps~~ ~~It is a maxheap~~

$A[1 \dots A.heapsize]$ is a maxheap.

In other words $A[1 \dots i]$ is a maxheap.

4.2

Priority Queues

A popular application of heaps is in priority queues. For example, when you want to schedule jobs on a shared machine, you want to run them according to their priority so you keep the job ~~values~~ in a maxheap and then run the job at the root and rearrange the heap accordingly.

In general a priority queue is a data structure for maintaining a set S of elements, each with an associated ~~key~~ value called a key. (e.g. $A[i]$ is the value/key of element i). A max-priority queue supports the following operations:

Insert(S, x): insert the element x to S (i.e. $S \leftarrow S \cup \{x\}$)

Maximum(S): return the element with maximum key in S .

Extract-Max(S): removes and returns the element in S with largest key

Increase-Key(S, x, k): increase the value/key of x to k

~~For~~ We will implement a priority queue using heaps. Applications would require that we store at each $A[i]$ also the "name of the element in S " when $S \neq \{1 \dots n\}$ (each pointer to relevant object etc.), but as this depends on the application we will ignore it.

So let us see how a maxheap can be used to support max-priority queue.

Heap-Maximum(A)

1. Return $A[1]$

4.8

Heap-Extract-Max(A)

1. If $A.\text{heap-size} < 1$
2. return "error". stop.
3. $\text{max} = A[i]$
4. $A[i] = A[A.\text{heap-size}]$
5. $A.\text{heap-size} = A.\text{heap-size} - 1$
6. Max-heapify(A, i)
7. Return max

Heap-Increase-Key(A, i, key)

1. If $\text{key} < A[i]$
2. return "error". stop.
3. $A[i] = \text{key}$
4. while $i > 1$ and $A[\text{Parent}(i)] < A[i]$
exchange $A[i]$ with $A[\text{Parent}(i)]$
5. $i = \text{Parent}(i)$

Max-Heap-Insert(A, key)

1. $A.\text{heap-size} = A.\text{heap-size} + 1$
2. $A[A.\text{heap-size}] = -\infty$
3. Heap-Increase-Key(A, A.heap-size, key)

Claim: algs. work and require $\mathcal{O}(\lg n)$ time.